

The architecture of a Recurrent Neural Network (RNN) is specifically designed for sequential data, such as time-series, text, or audio. RNNs are unique in that they use feedback loops to process sequences, allowing them to retain information from previous time steps and thus learn dependencies over time. Here's a breakdown of the key components and structure of an RNN:

1. Basic Components of an RNN Layer

- **Input Layer:** Receives sequential input data. For example, if processing a sentence, each word is processed in sequence, one at a time.
- **Recurrent (Hidden) Layer:** This layer has connections looping back onto itself, creating a "memory" that stores information from previous steps in the sequence.
 - The hidden layer at each time step has an associated hidden state, denoted by h_t .
 - At each time step t , the hidden state h_t is updated based on the current input x_t and the previous hidden state h_{t-1} .
- **Output Layer:** Produces the output for each time step, based on the hidden state.

2. Mathematical Representation of RNN Computation

- At each time step t , the RNN computes the hidden state h_t and output o_t as follows:

- **Hidden State Update:**

$$h_t = \tanh(W_h \cdot x_t + U_h \cdot h_{t-1} + b_h)$$

Here:

- W_h : Weight matrix between the input and the hidden layer.
- U_h : Weight matrix for the recurrent connection in the hidden layer.
- b_h : Bias for the hidden layer.
- h_{t-1} : Previous hidden state (memory from the last time step).
- x_t : Current input.
- \tanh : Activation function that introduces non-linearity.

- **Output Calculation:**

$$o_t = \text{softmax}(W_o \cdot h_t + b_o)$$

Here:

- W_o : Weight matrix between the hidden layer and the output layer.
- b_o : Bias for the output layer.
- o_t : Output at time t , often a probability distribution if working with classification.

3. Feedback Loop and Information Flow

- Each hidden state h_t in an RNN depends not only on the current input x_t but also on the hidden state from the previous time step h_{t-1} , creating a chain-like dependency across time.
- This feedback loop enables the RNN to retain information about earlier steps in the sequence, allowing it to model temporal dependencies and sequential patterns.

4. Unrolling the RNN

- During backpropagation, RNNs are often "unrolled" across the entire sequence to calculate gradients for each time step.
- This unrolled representation makes it easier to visualize the time dependencies, with each "unrolled" copy representing the network's state at a specific time step.

5. Limitations of RNNs

- **Vanishing/Exploding Gradients:** RNNs struggle with learning long-term dependencies due to vanishing or exploding gradients during backpropagation through time (BPTT).
- **Short-Term Memory:** Standard RNNs are typically effective for short sequences but face limitations when modeling long-term dependencies, as older information may be overwritten by new input.